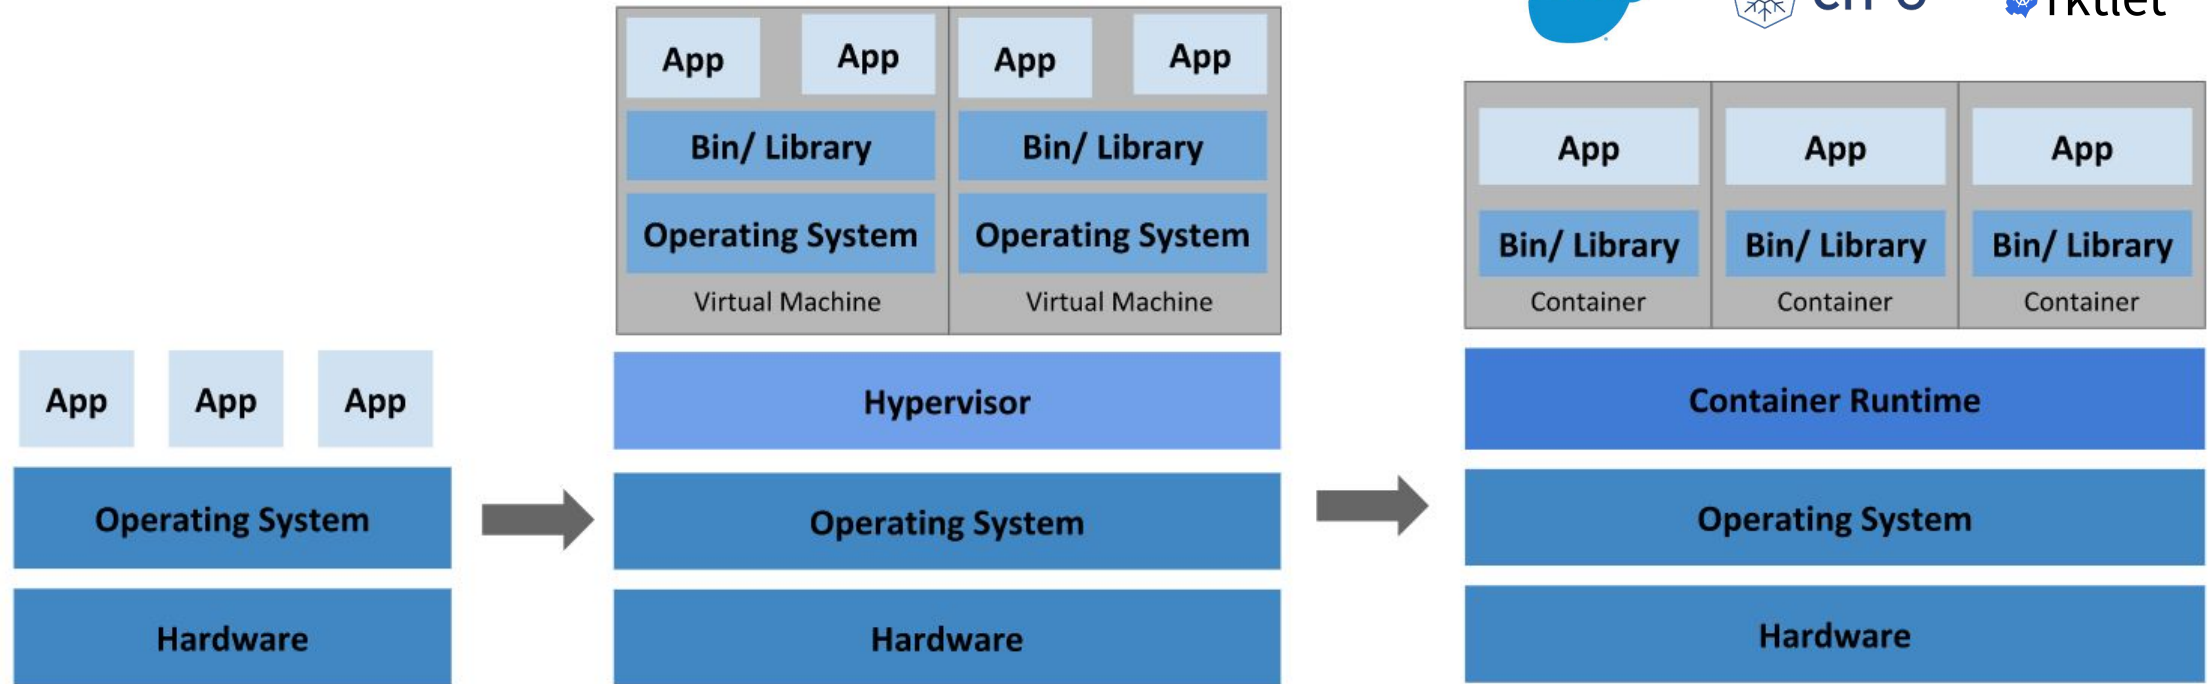


# Kubernetes for Developers

John Bush & Pre Bhakta  
Solutions Engineers  
VMware, Inc.

April 27, 2020

# Back To The Future ...



## Traditional Deployment

- Organizations ran applications on physical servers
- No way to define resource boundaries
- Not scalable
- Problematic in many ways

## Virtualized Deployment

- Run multiple Virtual Machines (VMs) on a single physical server's CPU
- Allows applications to be isolated between VMs
- Need Hypervisor to run the VM
- Provides a level of security as the information of one application cannot be freely accessed by another application.
- Each VM is a full machine running all the components, including its own operating system, on top of the virtualized hardware - Heavyweight (GB in size)

## Container Deployment

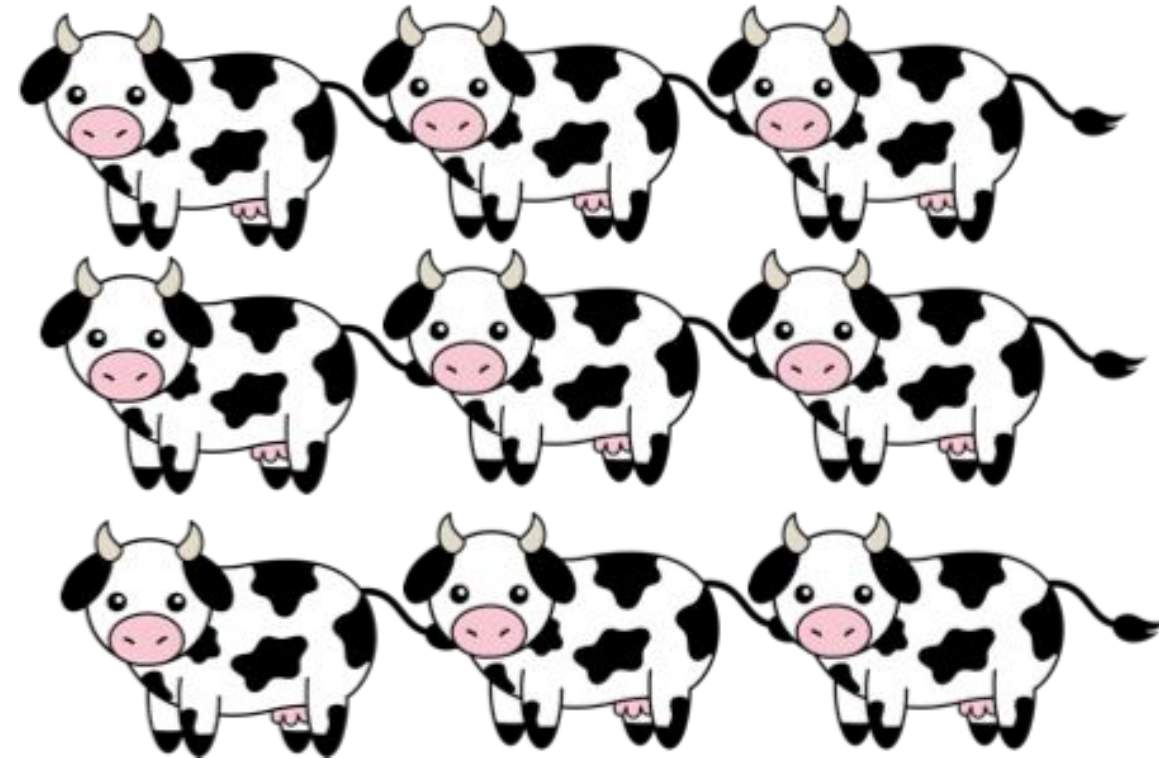
- Similar to VMs, but they have relaxed isolation properties to **share** the Operating System (OS) among the applications - lightweight (MB in size)
- Has its own filesystem, CPU, memory, process space, and more.
- Portable across clouds and OS distributions.
- Increased ease and efficiency of container image creation compared to VM image use.
- **Raises the level of abstraction from running an OS on virtual hardware TO running an application on an OS**
- Great way to bundle and run applications

# Think Of VMs vs. Containers As “Pets Vs. Cattle”



## Virtual Machine

Long lived  
You name them  
When they get ill, your nurse them back to health



## Container

Ephemeral  
You brand them with #'s  
When they get ill, you get another one

# Things Container Technologies Can't (Or Won't) Do ...



- Solve Port Mapping Hell
- Monitor Running Containers
- Handle Dead Containers
- Move Containers to Improve Utilization
- Autoscale Container Instances To Handle Load

kubernetes (n.) - *greek word for pilot or helmsman*



*Also called “K8s”*





# Back To The Cattle vs. Pets Thing ...



And you can never get away from  
“Pets” ....

**UNLESS**

You have an environment to  
manage the “Cattle” ...



# So what IS kubernetes?



A Production-Grade **Container Orchestration** System

Originally designed by **Google** (2014) and **donated** to the Cloud Native Computing Foundation (2015).

**Portable** and **extensible** open-source platform for managing containerized workloads.

It aims to provide a platform for **automating deployment, scaling, and operations** of application containers across **clusters** of hosts.

Principle: Manage your applications like **Cattle** instead of like Pets




# From Google Trends - Searches On “Kubernetes”

● **kubernetes**  
Search term

+ Compare

Worldwide ▾ 1/1/11 - 11/7/19 ▾ All categories ▾ Web Search ▾

Interest over time ?



# So Really ..... Why Is Kubernetes So Popular?

It's so much more than just enabling a containerized application to scale.

Kubernetes has release-management features that enable:

- It's vendor agnostic

- Updates with near-zero downtime

- Version rollback

- Clusters that can 'self-heal' when there is a problem

- Load balancing, auto-scaling and SSL can easily be implemented.

- Helm, a plugin for Kubernetes, is a package manager that makes deployment of apps and its dependencies very easy

- Remember the "Things that Container technologies can't / won't do?" slide ? .....

And it does all of the above with speed, power and efficiency !!!



# Every cloud supports Kubernetes



A re-architecture of the vSphere server virtualization platform that turns vSphere into a Kubernetes native platform.

# Want To Install Kubernetes?

Simplest

Most  
involved



Minikube



Google  
Container  
Engine  
(GKE)



AWS  
Provider



Manual  
install


# K8s Architecture Basics

---



# Nodes

- A Physical or Virtual Machine on which K8s is installed
- A node is a worker machine and where containers will be launched by K8s
- Also called a “minion” in the past
- Not advisable to have only one node









# Cluster

---

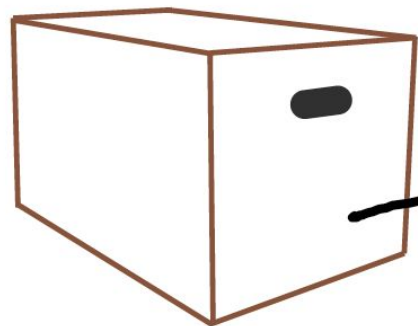


- Set of nodes grouped together
- Gives failover and ability to share the load

# Master



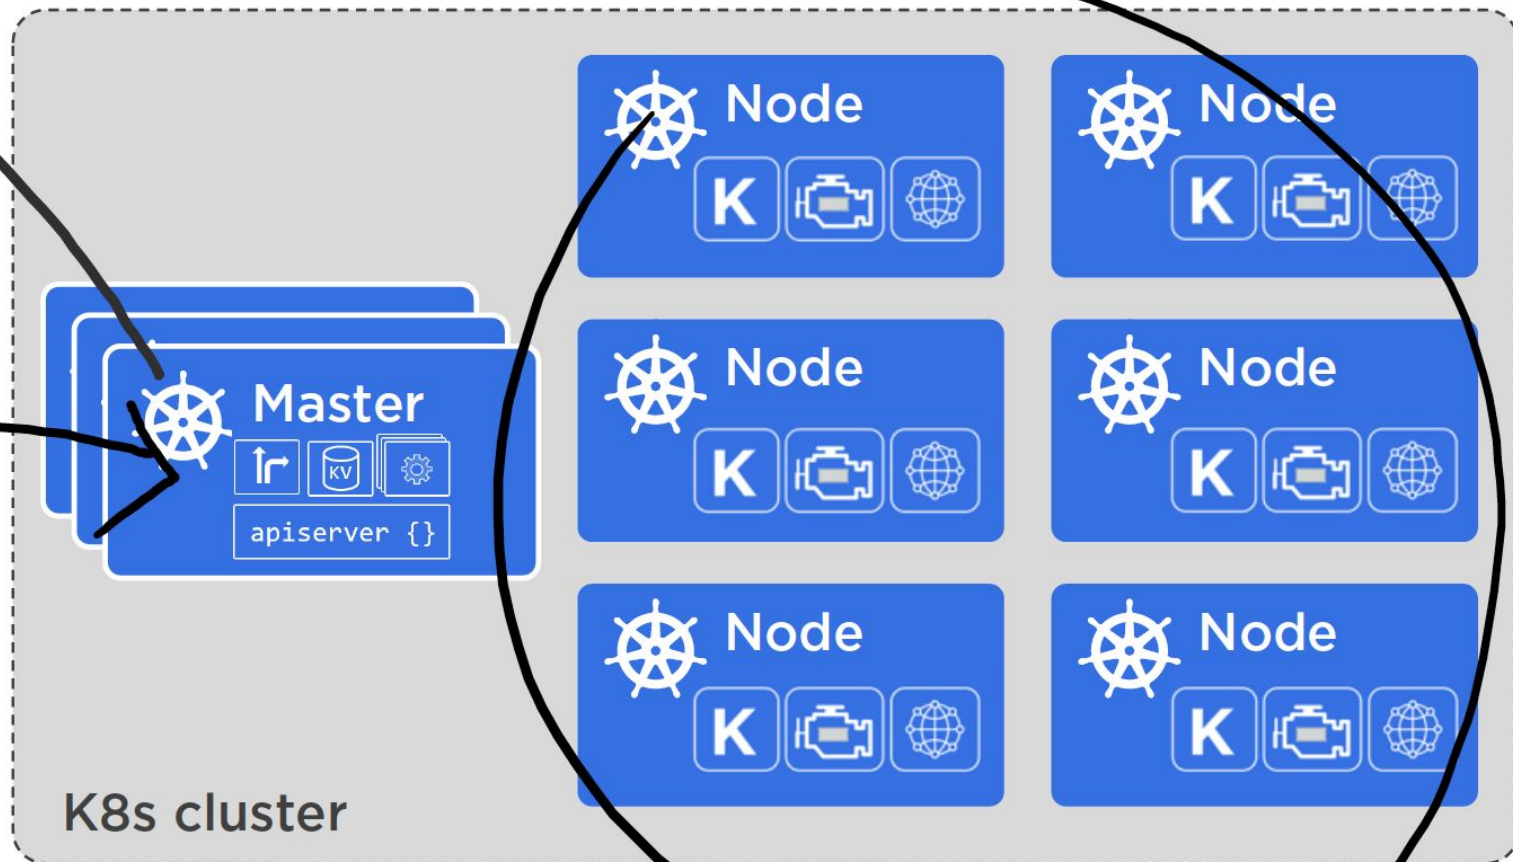
- Another node with K8s installed in it and configured to be a master node
- Master watches over the nodes in the cluster and is responsible for the actual orchestration of containers on the worker nodes
- Master Node:
  - Manages the cluster
  - Storage of the member cluster information
  - How are the nodes monitored?
  - When the node fails, how do you move the workload of the failed node to another worker node?



Deployment

*in charge!*


*do the work!*



Nodes a.k.a. Minions

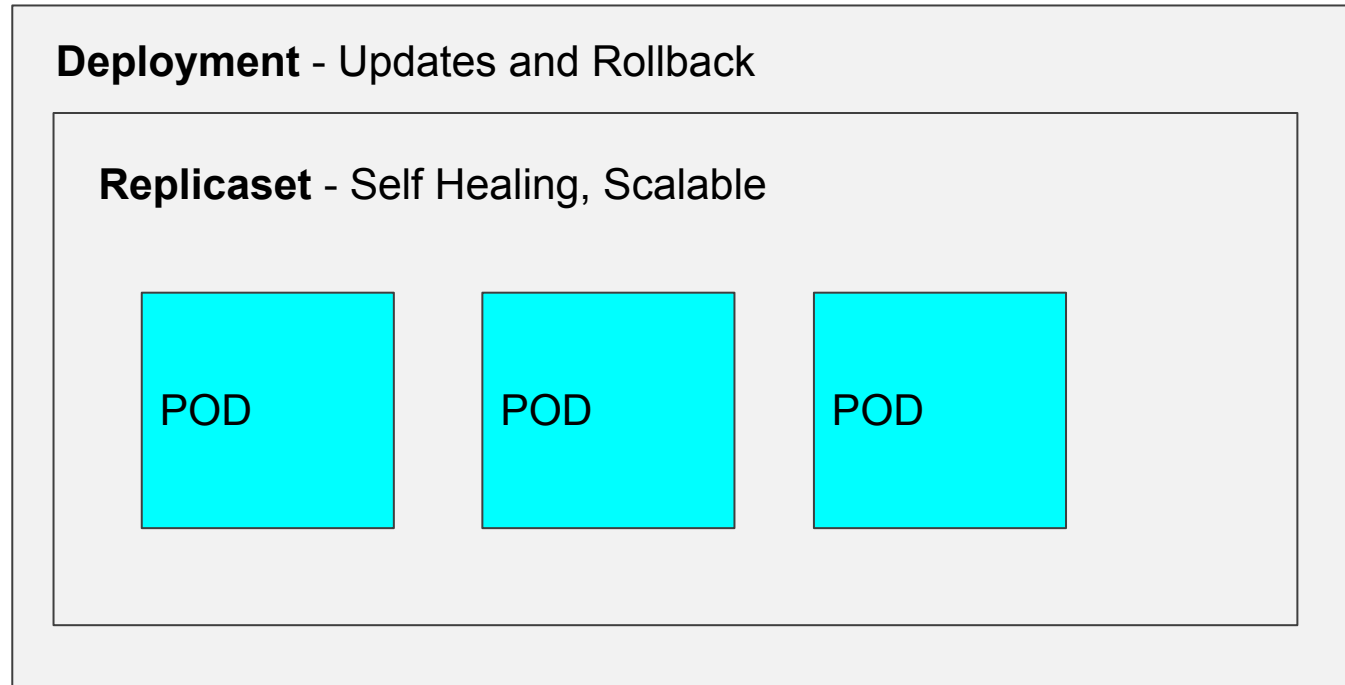


# POD



- K8s does not run containers directly; instead it wraps one or more containers into a higher level structure called a “POD”.
- PODs almost always have a 1:1 relationship with containers running an application
- A POD is the smallest object you can create in Kubernetes
- Any containers in the same POD will share the same resources and local network
- Each POD has a unique IP Address

# Replicaset, Deployments, & Services



**Service** - Enables network access to a set of pods

- ClusterIP
- NodePort
- LoadBalancer

# Configuring Kubernetes Using YAML Files ...

- “Yet Another Markup Language”
- Easy way to configure and manage K8 cluster(s)
- Using YAML for K8s definitions gives you a number of advantages, including:
  - Convenience:** You’ll no longer have to add all of your parameters to the command line
  - Maintenance:** YAML files can be added to source control, so you can track changes
  - Flexibility:** You’ll be able to create much more complex structures using YAML than you can on the command line
- YAML is a superset of JSON, which means that any valid JSON file is also a valid YAML file

```
---
apiVersion: v1
kind: Pod
metadata:
  name: rss-site
  labels:
    app: web
spec:
  containers:
    - name: front-end
      image: nginx
      ports:
        - containerPort: 80
    - name: rss-reader
      image: nickchase/rss-php-nginx:v1
      ports:
        - containerPort: 88
```




# Components

- \* Front end for Kubernetes
- \* Users, command line interfaces, management devices all talk to the API Server to interact with the K8 Cluster

- \* Looks for newly created Pods that have no Node assigned
- \* For every Pod that the scheduler discovers, the scheduler becomes responsible for finding the best Node for the Pod to run on

- \* The brains behind orchestration
- \* Watches the state of the cluster, then makes or request changes where needed

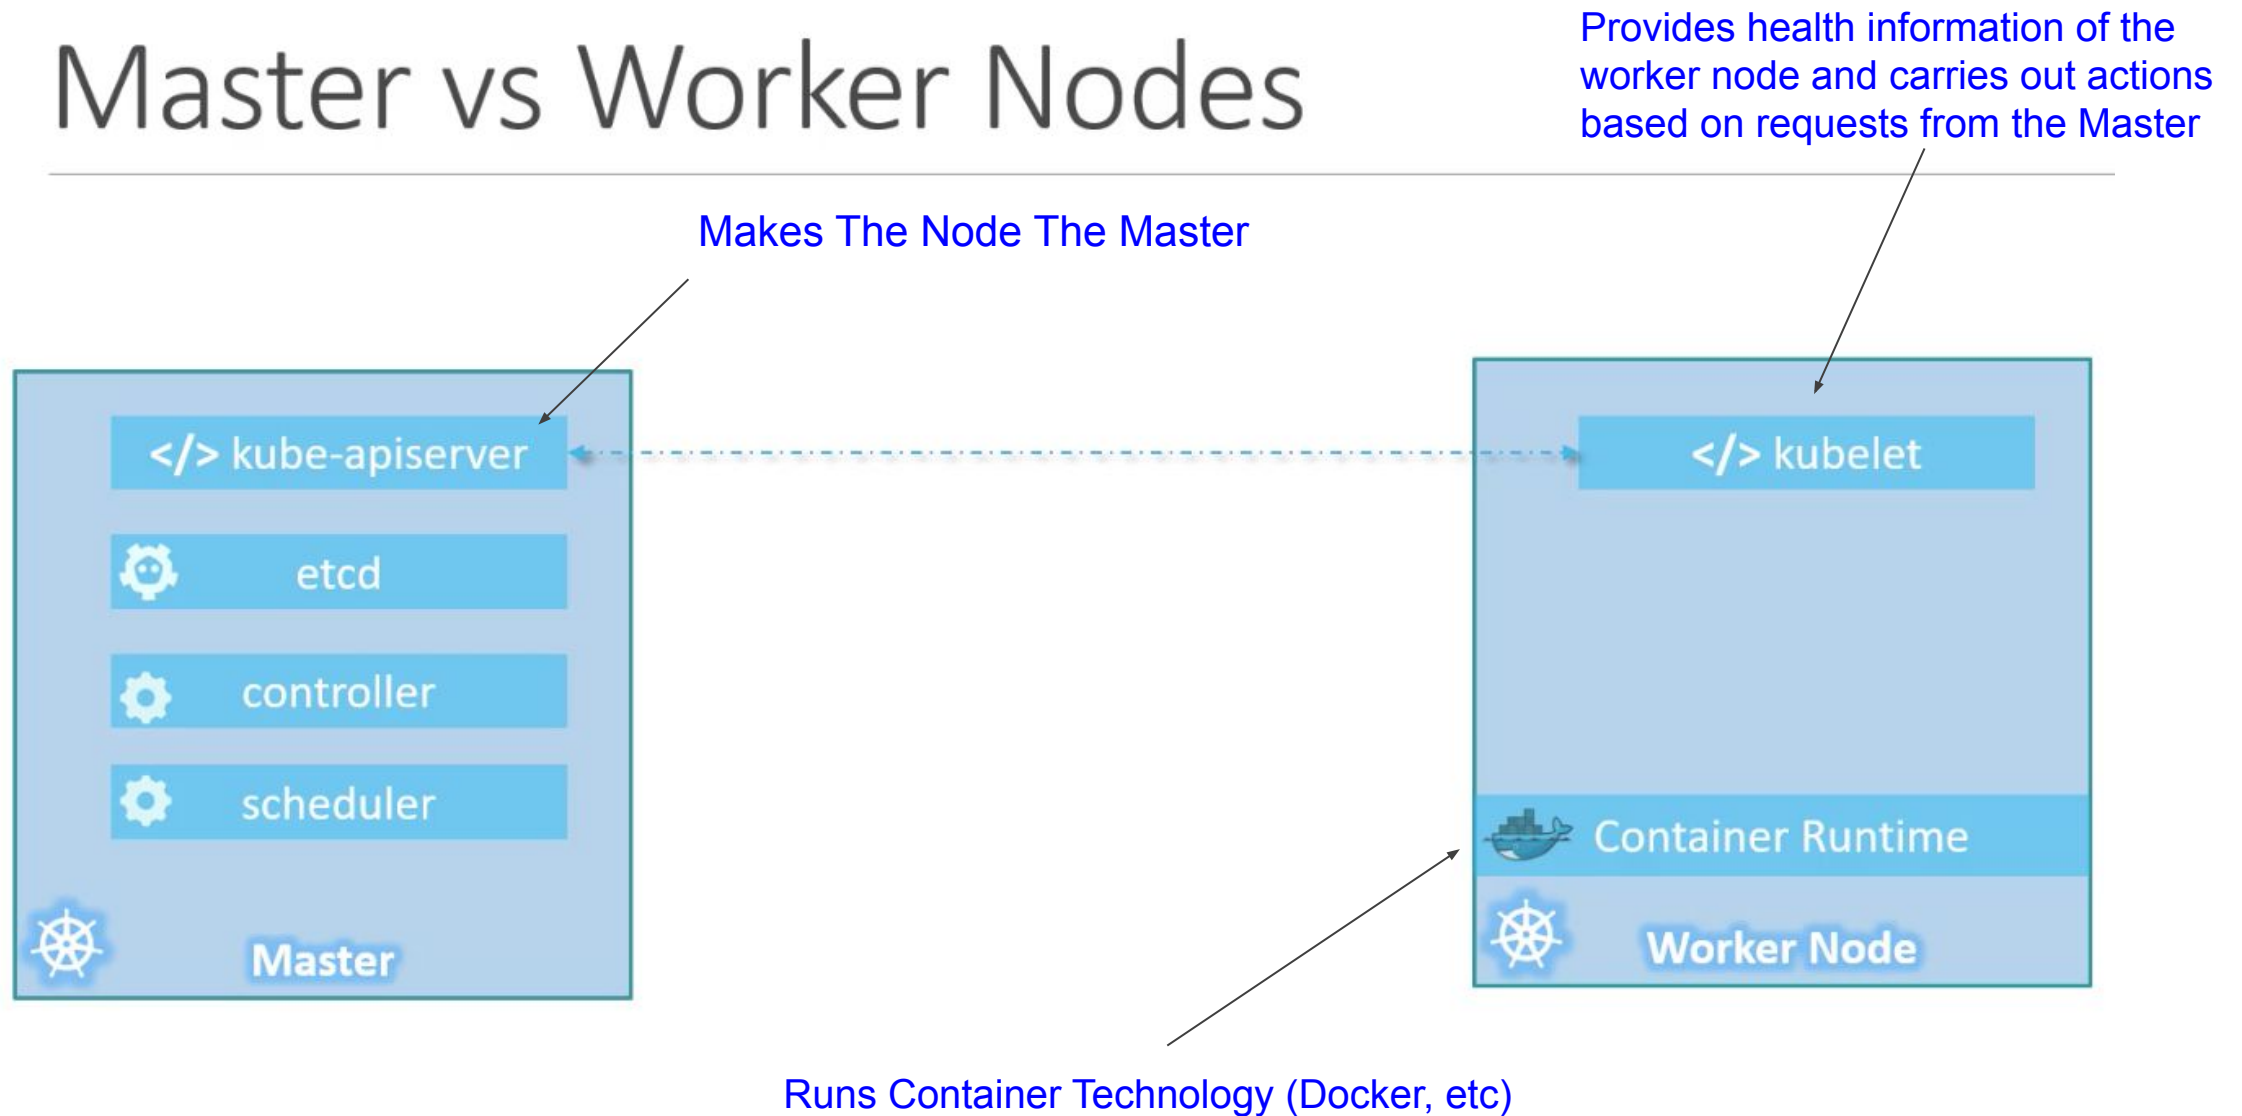


- \* Distributed key / value store used by K8 to store all data about the cluster
- \* K8s uses etcd to store all its data – its configuration data, its state, and its metadata

- \* Primary “node agent” that runs on each node of the cluster
- \* Agent makes sure that the containers running on the node are running as expected

- \* Underlying software that is used to run the containers
- \* In most cases, its Docker but there are other options

# Master vs Worker Nodes



# kubectl

- command line interface for running commands against Kubernetes clusters

---

```
kubectl run hello-minikube
```



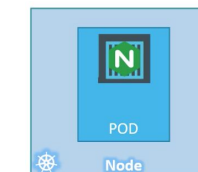
```
kubectl cluster-info
```



```
kubectl get nodes
```



```
kubectl run nginx --image nginx
```





# Demo Time

# Thank You

<https://github.com/fjb4/kubernetes-demo>